

# Lab Notes

## CMPU4021 Distributed Systems

### ***Concurrency and Multithreading in Java***

# Concurrency

- Concurrency in distributed systems
  - Concurrent requests to its resources
  - Each resource must be designed to be safe in a concurrent environment
- Concurrent programming
  - Systems can do more than one thing at a time
  - E.g. streaming audio application must simultaneously read the digital audio off the network, decompress it, manage playback, and update its display.
  - The word processor should always be ready to respond to keyboard and mouse events, no matter how busy it is reformatting text or updating the display.
  - Software that can do such things is known as *concurrent* software.
- In concurrent programming, there are two basic units of execution:
  - *processes* and *threads*
- `java.util.concurrent` *packages*

# Processes

- A process has a self-contained execution environment
- Has a complete, private set of basic run-time resources
- Each process has its own memory space
- To facilitate communication between processes, most operating systems support *Inter Process Communication* (IPC) resources
  - such as pipes and sockets
- IPC is used not just for communication between processes on the same system, but processes on different systems

# Multiprocess applications

- Most implementations of the Java virtual machine run as a *single* process.
- A Java application can create additional processes using a Most implementations of the Java virtual machine run as a single process.
- A Java application can create additional processes using a `ProcessBuilder` object.

# Threads

- Sometimes called *lightweight processes*
- A **thread** is a single sequential flow of execution that runs through a program.
- Threads exist within a process
  - every process has at least one
- Unlike a process, a thread does not have a separate allocation of memory, but shares memory with other threads created by the same application.

# Threads

- Threads share the process's resources, including memory and open files.
  - This makes for efficient, but potentially problematic, communication
- Multithreading
  - You can have more than one thread running at the same time inside a single program, which means it shares memory with other threads created by the same application.
- Every application has at least one thread (started with `main()`) — or several, if you count "system" threads that do things like memory management and signal handling.

# Why use threads?

- Threads can help in creating better applications, e.g.:
  - print in the ‘background’ (while editing);
  - scrolling through web pages texts while the browser is busy fetching the images.
- Threads can quicken calculations
  - having different threads executing sub-tasks

# Concurrent execution and context switching

- Two types of phases:
  - Compute bound phases - CPU is utilised for various calculations.
  - I/O bound phases – require the various input and output devices (printers, hard-disks, network cards etc.), CPU is mostly idle, waiting for the I/O device to do its task.
- Interleaving of phases.
- *Context switch*: the process when one thread relinquish the CPU and another thread starts running it.
- Thread *context*: each running thread has its own point of execution and private view of the values of local variables.



# Why not use threads?

- Context switch is costly
- Extra CPU tasks: 'freeze' and 'de-freeze' state of threads
- Single CPU multithreading can take more time than having linear code to return the same result

# Threading with Java

- At least one thread
  - launched by JVM when `main` is started and ‘killed’ when `main` terminates.
- *Multithreading*
  - using one or more extra threads in order to ‘offload’ processing tasks onto them.
  - These threads need to be programmed explicitly.

# Single processor threading

- Threads with the same priority are each given an equal **time-slice** or **time quantum** for execution on the processor;
- **Pre-emption** (more urgent attention) – if thread is assigned higher priority.

# Using threads in Java

- Multithreading directly accessible – no need to go through an operating system API
- System independent
- There are two basic strategies for using Thread objects to create a concurrent application.
  - To directly control thread creation and management, simply *instantiate* Thread each time the application needs to initiate an asynchronous task.
  - To abstract thread management from the rest of your application, pass the application's tasks to an *executor*.
- We will concentrate on the first way

# Defining and Starting a Thread

Two ways for creating threads:

- *Provide a Runnable object*
  - Create a class that does not extend `Thread` and specify explicitly that it implements interface `Runnable`.
  - The `Runnable` interface defines a single method, `run`, meant to contain the code executed in the thread. The `Runnable` object is passed to the `Thread` constructor
- *Subclass Thread*
  - The `Thread` class itself implements `Runnable`, though its `run` method does nothing.
  - Create a class that extends `java.lang.Thread` class

## Threading – Way 1: *Provide a Runnable object*

```
public class HelloRunnable implements Runnable {  
  
    public void run() {  
        System.out.println("Hello from a thread!");  
    }  
  
    public static void main(String args[]) {  
        (new Thread(new HelloRunnable())).start();  
    }  
  
}
```

# *Threading - Way 2: Subclass Thread*

```
public class HelloThread extends Thread {  
  
    public void run() {  
        System.out.println("Hello from a thread!");  
    }  
  
    public static void main(String args[]) {  
        (new HelloThread()).start();  
    }  
  
}
```

## Extend the `Thread` or implement the `Runnable`?

- If your class *must* be derived from some other class (e.g. `Applet`) then it should implement `Runnable`, otherwise it should extend `Thread`.
- If the class has a superclass (other than `Object`), it can not extend the `Thread` (no multiple inheritance). A class that implements the `Runnable` interface can extend *any* class.



# Extending the `java.lang.Thread` class

- The class that extends `Thread` should override the `run` method of class `Thread`.
- The `run` method serves the same purpose as the method `main` for a full application. Like `main`, `run` may not be called directly; it specifies the actions that a thread is to execute.
- The most common constructors used:
  - `Thread()` - allocates a new `Thread` object; the system generates a name of the thread as: *Thread-n*
  - `Thread(String name)` - allocates a new `Thread` object with the name *name*.

# The Thread class (I)

- **getName()** method – retrieves the name of thread. E.g.:

```
Thread firstThread = new Thread();  
Thread secondThread() = new Thread("namedThread");  
System.out.println(firstThread.getName());  
System.out.println(secondThread.getName());
```

**output:**

```
Thread-0  
namedThread
```

# Pausing Execution with Sleep

- `Thread.sleep()` causes the current thread to suspend execution for a specified period
- This is an efficient means of making processor time available to the other threads of an application or other applications that might be running on a computer system

# Pausing Execution with Sleep

- Method

```
public static void sleep(long millis)  
                        throws InterruptedException
```

causes the currently executing thread to sleep (temporarily cease execution) for the specified number of milliseconds.

```
firstThread.sleep(1500); // pause for 1.5 seconds
```

- It allows other threads to be executed. When the sleeping time expires, the sleeping thread returns to a *ready* state, waiting for the processor.
- Sleep time very often determined by using the `java.lang.Math.random()` method in order to achieve a random sleeping time for threads.

# Interrupts

- An interrupt is an indication to a thread that it should stop what it is doing and do something else.
- It's up to the programmer to decide exactly how a thread responds to an interrupt, but it is very common for the thread to terminate.
- 
- A thread sends an interrupt by invoking `interrupt` on the `Thread` object for the thread to be interrupted.

# Interrupts

- The method:  
`public void interrupt()`  
Interrupts an individual thread.
- It may be used by other threads to ‘awaken’ a sleeping thread before the thread’s sleeping time has expired. Used very rarely.
- In the next example, static method `random` is used to generate a random sleeping time for each of two threads that display their own names 10 times.

# Example

```
public class ThreadShowName extends Thread {
    public static void main(String[] args){
        ThreadShowName thread1, thread2;
        thread1 = new ThreadShowName();
        thread2 = new ThreadShowName();

        thread1.start();
        thread2.start();
    }

    public void run(){
        int pause;
        for (int i=0; i<10; i++){
            try{
                System.out.println(getName() + " being executed.");
                pause = (int)(Math.random() * 3000);
                sleep(pause);
            }
            catch (InterruptedException e){
                System.out.println(e.toString());
            }
        }
    }
}
```

# Joins

- The join method allows one thread to wait for the completion of another, so that a thread will not start running until another thread has ended.
- If myThread is a Thread object whose thread is *currently* executing,

```
myThread.join();
```

- Causes the current thread to pause execution until myThread's thread terminates.
- Overloads of join allow the programmer to specify a waiting period.
- 
- As with sleep, join is dependent on the OS for timing, so you should not assume that join will wait exactly as long as you specify.
- Like sleep, join responds to an interrupt by exiting with an InterruptedException.



# Joins: Example

```
threadMessage("Waiting for MessageLoop thread to finish");
// loop until MessageLoop thread exits
while (t.isAlive()) {
    threadMessage("Still waiting...");
    // Wait maximum of 1 second for MessageLoop thread to finish.
    t.join(1000);
    if (((System.currentTimeMillis() - startTime) > patience)
        && t.isAlive()) {
        threadMessage("Tired of waiting!");
        t.interrupt();
        // Shouldn't be long now -- wait indefinitely
        t.join();
    }
}
threadMessage("Finally!");
```

# Implementing the Runnable interface

1. Create an application class that explicitly implements the `Runnable` interface.
2. In order to create a thread – instantiate an object of the `Runnable` class and ‘wrap’ it in a `Thread` object (by creating a `Thread` object and passing the `Runnable` object as an argument to the `Thread` constructor). E.g.

```
Thread(Runnable<object>)
```

```
Thread(Runnable<object>, String<name>)
```

When either of these constructors is used, The `Thread` object uses the `run` method of the `Runnable` object in place of its own (empty) `run` method.

# Implementing the Runnable interface (I)

```
class PrimeRun implements Runnable {
    long minPrime;
    PrimeRun(long minPrime) {
        this.minPrime = minPrime;
    }
    public void run() {
        // compute primes larger than minPrime
        . . .
    }
}
```

The following code would then create a thread and start it running:

```
PrimeRun p = new PrimeRun(143);
Thread thread1 = new Thread(p);
thread1.start();
```

**Or, shorter:**

```
Thread thread1 = new Thread(new PrimeRun(143));
thread1.start();
```

# Extending the Thread / Implementing Runnable

```
class MyThread extends Thread {  
    public void run() {  
        System.out.println ("Hello World!");  
    }  
}
```

```
class MyTest {  
    public static void main(String Args[]) {  
        new MyThread().start();  
    }  
}
```

---

---

```
class MyThread implements Runnable {  
    public void run() {  
        System.out.println ("Hello World!");  
    }  
}
```

```
class MyTest {  
    public static void main(String Args[]) {  
        new Thread(new MyThread()).start();  
    }  
}
```

# Multithreaded servers

- Servers can handle more than one client/connection at a time.
- Two stages:
  - the main thread (running in method `main`) allocates individual threads to incoming clients.
  - the thread allocated to each individual client then handles all subsequent interaction between that client and the server (via the thread's `run` method).

# Multithreaded servers

- For each client-handling thread that is created, the main thread must ensure that the *client-handling thread is passed a reference to the socket* that was opened for the associated client.

(See the *MultiEchoServer.java* in Labs)

# Daemon Threads

- Daemon threads: that exist for the purpose of providing a certain service.
  - E.g. A daemon thread, named Background Image Reader, of the HotJava web browser, reads images from the file system or the network for any object or thread that needs an image.
- The **run()** method for a daemon thread is usually an infinite loop that waits for a service request.
- **setDaemon()** – any thread can become a daemon or return from a daemon mode to a non-daemon.
- **isDaemon()**

# Daemon Threads

- Normal thread and daemon threads differ in what happens when they exit.
- When the JVM halts any remaining daemon threads are *abandoned*: finally blocks are not executed, stacks are not unwound – JVM just exits.
  - Hence, daemon threads should be used sparingly and it is dangerous to use them for tasks that might perform any sort of I/O.



# Parallel Execution of Threads

```
class PrintThread implements Runnable {
    String str;
    public PrintThread (String str) {
        this.str = str;
    }
    public void run() {
        for (;;) System.out.print (str);
    }
}

class ConcurrencyTest {
    public static void main (String Args[]) {
        new Thread(new PrintThread("A")).start();
        new Thread(new PrintThread("B")).start();
    }
}
```

# Parallel Execution of Threads

- The output of the program above should look something like this (on multi-processor machines):

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAABBBBBBBBBBBBBBBBBBBBBBBBBBBB
BBBBBBBBBBBBBBBBBAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAABBBB
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAABBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAABBBBBBBBBBBBBBBBBBBBBBBB
```

- It has nearly equal number of A's and B's.:

# Preemptive versus Non-Preemptive multi-threading

- *Preemptive multi-threading* means that a thread may be preempted by another thread with an equal priority while it is running.
- The Java runtime will not preempt the currently running thread for another thread of the same priority. However, the underlying operating system implementation of threads may support preemption.
- Today, nearly all operating systems support preemptive multitasking, including the current versions of Windows, Mac OS, GNU/Linux, iOS and Android.
- If multi-tasking is not preemptive, the output of the previous example program (on an older OS) would like something like this:

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAA
```

# Well behaved Threads

- A thread is supposed to be well behaved and give up the CPU periodically in order for other threads to be able to run.
- If your thread does not give up the CPU by suspending itself, waiting for a condition, sleeping or doing I/O operations then it should relinquish the CPU periodically by invoking the *Thread* class's *yield()* method.

# Preemptive versus Non-Preemptive multi-threading

```
class WellBehavedPrintThread implements Runnable {  
    String str;  
    public PrintThread (String str) {  
        this.str = str;  
    }  
    public void run() {  
        for (;;) {  
            System.out.print (str);  
            Thread.currentThread().yield();  
        }  
    }  
}
```

# Preemptive versus

# Non-Preemptive multi-threading

- The statement `Thread.currentThread().yield()` uses a public static method of the *Thread* class to get a handle to the currently running thread and then tells it to yield.
- The output of this example is:

[illegible]

- As a rule of thumb, threads should yield whenever possible, to allow others run.

# Thread States

A thread can be in any of the following states:

- **NEW**
  - A thread that has not yet started is in this state.
- **RUNNABLE**
  - A thread executing in the Java virtual machine is in this state.
- **BLOCKED**
  - A thread that is blocked waiting for a monitor lock is in this state.
- **WAITING**
  - A thread that is waiting indefinitely for another thread to perform a particular action is in this state.
- **TIMED\_WAITING**
  - A thread that is waiting for another thread to perform an action for up to a specified waiting time is in this state.
- **TERMINATED**
  - A thread that has exited is in this state.

# Thread Priorities

- Each thread is assigned a priority, ranging from MIN\_PRIORITY (equals 1) to MAX\_PRIORITY (which is 10). A thread inherits its priority from the thread that spawned it.
- The scheduling algorithm always lets the highest priority runnable thread run. If at any time a thread with a higher priority than all other runnable threads becomes runnable, the runtime system schedules it for execution.
- The new higher priority thread is said to preempt the other threads.
- A lower priority thread can only run when all higher priority threads are non-runnable.



# Thread Priorities

- The priority of the main thread (the thread that starts the `main()` method of your program) is `NORM_PRIORITY` (equals 5).
- You can set a thread's priority by invoking the `setPriority()` method and get the thread's priority by invoking `getPriority()`.

# Thread synchronization

- Many threads - need to synchronize their activities.
- Need to prevent concurrent access to data structures in the program that are shared among the threads.
- Java provides mechanisms for synchronization and mutual exclusion (allowing only one thread to run through critical code sections in the program).

# Thread Synchronization

- Threads communicate primarily by sharing access to fields and the objects reference fields refer to.
- Two kinds of errors possible:
  - thread interference; and
  - memory consistency errors.
- The tool needed to prevent these errors is synchronization.

# Thread Synchronization

- Synchronization can introduce thread contention, which occurs when two or more threads try to access the same resource simultaneously and cause the Java runtime to execute one or more threads more slowly, or even suspend their execution.
- Starvation and livelock are forms of thread contention.

# Starvation

- The term *starvation* is used to denote situations where one thread is deprived of a resource (namely an access to a monitor). Unlike deadlock, in a starvation situation the calculation can continue in the system, it's just the starved thread that can't go on.
- Starvation can occur when a high priority thread starts running and never relinquish the CPU.
- Although the high priority thread can do great many things, all the lower priority threads are starved.

# Livelock

- A thread often acts in response to the action of another thread. If the other thread's action is also a response to the action of another thread, then livelock may result.
- Livelocked threads are unable to make further progress.
  - However, the threads are not blocked — they are simply too busy responding to each other to resume work.
  - This is comparable to two people attempting to pass each other in a corridor:
    - John moves to his left to let Paul pass, while Paul moves to his right to let John pass.
    - Seeing that they are still blocking each other, John moves to his right, while Paul moves to his left. They're still blocking each other, so...

# Deadlocks

- *Deadlock* – occurs when threads are waiting for events that will never occur
- Deadlock describes a situation where two or more threads are blocked forever, waiting for each other.
- The two threads are in a deadlock if one of them waits for the value to change while blocking the other one from changing it.
- Deadlock occurs when two or more threads are waiting for some condition to change, while that condition is precluded from changing because of all threads that can change the condition are waiting.

# Thread synchronization (I)

MyData (Version 1):

```
class MyData {  
    private int Data;  
    public void store(int Data) {  
        this.Data=Data;  
    }  
    public int load() {  
        return this.Data;  
    }  
}
```

- supplies a store/load interface:
- stores integer data in a private property,



# Producer/Consumer Problem

Example:

- Two threads:
  - The first is trying to store values (producer)
  - and the second is trying to fetch the stored values (consumer)

```
class Main { // This class is used to set things in
    motion.
    public static void main(String argv[]) {
        MyData data = new MyData();
        new Thread(new Producer(data)).start();
        new Thread(new Consumer(data)).start();
    }
}
```

# Producer

```
class Producer implements Runnable {
    MyData data;
    public Producer(MyData data) {
        this.data = data;
    }
    public void run() {
        int i;
        for (i=0;;i++) {
            data.store(i);
            System.out.println ("Producer: "+i);
            try {
                // doze off for a random time (0 to 0.5 sec)
                Thread.sleep ((int) (Math.random()*500));
            } catch (InterruptedException e) { }
        }
    }
}
```

# Consumer

```
class Consumer implements Runnable { // The consumer
    MyData data;
    public Consumer(MyData data) {
        this.data=data;
    }
    public void run() {
        for (;;) {
            System.out.println ("Consumer: "+data.load());
            try {
                // sleep for a random time (0 to 0.5 sec)
                Thread.sleep ((int) (Math.random()*500));
            } catch (InterruptedException e) { }
        }
    }
}
```

# Problem

## Output:

Producer: 0

Consumer: 0

Producer: 1

Consumer: 1

Consumer: 1

Producer: 2

Producer: 3

Consumer: 3

Producer: 4

Producer: 5

- The producer has no way to know that the consumer did not consume the data yet, and therefore he overwrites it.
- The consumer has no way of knowing if the value it reads is a new value or the old one.

# Solution 1

- Use ordinary Boolean variables as flags to control data access:
  - The *Ready* flag will signify that new data was produced and is ready to be consumed;
  - The *Taken* flag will signify that the data was consumed and it is okay to overwrite it.
- Disadvantage:
  - the *store()* and *load()* methods are using *busy-loops*. The threads are constantly checking the flags to see if their value has been changed.

# Solution 1 – MyData (Version 2)

```
class MyData {  
    private int data;  
    private boolean ready;  
    private boolean taken;  
  
    public MyData() {  
        ready = false;  
        taken = true;  
    }  
    public void store(int data) {  
        while (!taken);  
  
        this.data = data;  
        taken = false;  
        ready = true;  
    }  
    public int load() {  
        int data;  
        while (!ready);  
        data = this.data; // save the value because after Taken turns "true" it may change at any time.  
        ready = false;  
        taken = true;  
        return data;  
    }  
}
```

# Intrinsic locks or monitor locks

- Synchronization is built around an internal entity known as the intrinsic lock or monitor lock.
  - The API specification often refers to this entity simply as a "monitor."
- A monitor is associated with a specific data item and functions as a lock on that data.
- When a thread holds the monitor for some data item, other threads are locked out and cannot inspect or modify the data.
- A thread can acquire the monitor if no other thread currently owns it, and it can release it at will.
- A thread can re-acquire the monitor if it already owns it.
- A locking is achieved by using the Java keyword `synchronized`
  - *synchronized methods* and *synchronized statements*.

# Synchronized methods

- Declaring a method **synchronized** means that only the thread holding the monitor can run through the synchronized method in that instance.

```
public synchronized void updateSum(int amount) {  
    sum+=amount;  
}
```

- If *sum* is not locked when the above method is invoked, then the lock on *sum* is obtained, preventing any other thread from executing *updateSum*.



# *Synchronized statements*

- Any code segment can be declared as synchronized:
  - performance issues (if too many)
- Unlike synchronized methods, synchronized statements must specify the object that provides the intrinsic lock:

```
public void addName(String name) {  
    synchronized(this) {  
        lastName = name;  
        nameCount++;  
    }  
    nameList.add(name);  
}
```

# Solution 2 – MyData (Version 3)

```
class MyData {  
    private int data;  
    private boolean ready;  
    private boolean taken;  
  
    public MyData() {  
        ready = false;  
        taken = true;  
    }  
  
    public synchronized void store(int data) {  
        while (!taken);  
        this.data = data;  
        taken = false;  
        ready = true;  
    }  
  
    public synchronized int load() {  
        while (!ready);  
        ready = false;  
        taken = true;  
        return this.data;  
    }  
}
```

# Solution 2 (cont)

- Declaring the methods synchroniz-ed removes the need for storing the value of the Data variable in the load() method. The load() and store() will not be able to execute at the same time in different threads anymore.
- Problem:
  - When one thread will engage the busy-wait loop, still owning the monitor. The other thread will never be able to execute its code because it can not acquire the monitor. The two threads are in a deadlock: one of them waits for the value to change while blocking the other one from changing it.

# Solution 3

- We need is to acquire the monitor *after* waiting for the flag.
- Therefore, we use the `synchronized` keyword on a code segment (as opposed to an entire method), thus protecting only the *critical code segment* that needs mutual-exclusion.
- When using the `synchronized` keyword on a code segment, an object to synchronize upon needs to be supplied. This is the object whose monitor is to be used for the critical section.

# Solution 3 (cont) – MyData (Version 4)

```
class MyData {  
    private int data;  
    private boolean ready;  
    private boolean taken;  
    public MyData() {  
        ready = false;  
        taken = true;  
    }  
    public void store(int data) {  
        while (!taken);  
        synchronized (this) {  
            this.data = data;  
            taken = false;  
            ready = true;  
        }  
    }  
    public int load() {  
        while (!ready);  
        synchronized (this) {  
            ready = false;  
            taken = true;  
            return this.data;  
        }  
    }  
}
```

# Solution 3 (cont)

## Problem:

- the busy-wait loop: It is considered bad practice for a thread to use a busy-wait loop: it is processor-expensive on a preemptive implementation, and might cause deadlock on a non-preemptive one.
- In order to improve thread efficiency and to help avoid deadlock, the following methods are used:
  - `java.lang.Object.wait()`
  - `java.lang.Object.notify()`
  - `java.lang.Object.notifyAll()`

# Waiting for events: `java.lang.Object` methods

- They can just be called only:
  - When the current thread has a lock on the object (i.e. from within a `synchronized` method, or
  - From within a method that has been called by a `synchronized` method.)
- `wait()`
  - If a thread executing a `synchronized` method determines that it cannot proceed then it may put itself into a waiting state by calling method `wait()`. This releases the thread's lock on the shared object and allows other threads to obtain the lock.
- `notify()`
  - A `synchronized` method reaches completion, then it may call `notify()`, which will 'wake up' a thread that is in the waiting state.
- `notifyAll()`
  - 'wakes up' all object waiting on a given object.

# wait()

- Causes the current thread to wait until either another thread invokes the notify() method or the notifyAll() method for this object, or a specified amount of time has elapsed.
- The current thread must own this object's monitor.
- This method causes the current thread (call it T) to place itself in the wait set for this object and then to relinquish any and all synchronization claims on this object.
- Thread T becomes disabled for thread scheduling purposes and lies dormant until one of four things happens:
  - Some other thread invokes the notify method for this object and thread T happens to be arbitrarily chosen as the thread to be awakened.
  - Some other thread invokes the notifyAll method for this object.
  - Some other thread interrupts thread T.
  - The specified amount of real time has elapsed, more or less. If timeout is zero, however, then real time is not taken into consideration and the thread simply waits until notified.
- The thread T is then removed from the wait set for this object and re-enabled for thread scheduling.
- It then competes in the usual manner with other threads for the right to synchronize on the object; once it has gained control of the object, all its synchronization claims on the object are restored to the status quo ante - that is, to the situation as of the time that the wait method was invoked.
- Thread T then returns from the invocation of the wait method.
- Thus, on return from the wait method, the synchronization state of the object and of thread T is exactly as it was when the wait method was invoked.



# notify()

- Wakes up a single thread that is waiting on this object's monitor. If any threads are waiting on this object, one of them is chosen to be awakened.
- The choice is arbitrary and occurs at the discretion of the implementation.
- A thread waits on an object's monitor by calling one of the wait methods.
- The awakened thread will not be able to proceed until the current thread relinquishes the lock on this object.
- The awakened thread will compete in the usual manner with any other threads that might be actively competing to synchronize on this object;
  - for example, the awakened thread enjoys no reliable privilege or disadvantage in being the next thread to lock this object.
- This method should only be called by a thread that is the owner of this object's monitor. A thread becomes the owner of the object's monitor in one of three ways:
  - By executing a synchronized instance method of that object.
  - By executing the body of a synchronized statement that synchronizes on the object.
  - For objects of type Class, by executing a synchronized static method of that class.

# notifyAll()

- Wakes up all threads that are waiting on this object's monitor.
- A thread waits on an object's monitor by calling one of the wait methods.
- The awakened threads will not be able to proceed until the current thread relinquishes the lock on this object.
- The awakened threads will compete in the usual manner with any other threads that might be actively competing to synchronize on this object;
  - for example, the awakened threads enjoy no reliable privilege or disadvantage in being the next thread to lock this object.
- This method should only be called by a thread that is the owner of this object's monitor.

# Final Producer/Consumer Solution

- The `wait()` method makes a thread release the monitor and shifts from the *runnable* state to the *non-runnable* state.
- The thread will stay in a non-runnable state until it is waken up by a call to `notify()`.
  - When a thread stops `wait()`ing, it re-acquires the monitor.
- The `notify()` method arbitrarily chooses a thread from those that are `wait()`ing and releases it from its wait state.

# Final Producer/Consumer Solution

```
class MyData {  
    private int Data;  
    private boolean Ready;  
    public MyData() {  
        Ready=false;  
    }  
    public synchronized void store(int Data) {  
        while (Ready)  
            try {  
                wait();  
            } catch (InterruptedException e) { }  
        this.Data=Data;  
        Ready=true;  
        notify();  
    }  
    public synchronized int load() {  
        while (!Ready)  
            try {  
                wait();  
            } catch (InterruptedException e) { }  
        Ready=false;  
        notify();  
        return this.Data;  
    }  
}
```

# Deadlock

- We've seen how a deadlock can occur if we use a busy-waiting loop inside a monitor, waiting for another thread to change a condition but never giving it the opportunity to obtain the monitor.
- We say that threads are locked in a deadlock if the calculation (or whatever it is the threads are doing) can't continue.

# Semaphores

- Semaphores are often used to restrict the number of threads that can access some (physical or logical) resource.
- Each `acquire()` blocks if necessary until a permit is available, and then takes it.
- Each `release()` adds a permit, potentially releasing a blocking acquirer.

# java.util.concurrent.Semaphore

- Semaphore object maintains a set of permits:

## Constructors:

*Semaphore(int permits);*

- Each acquire blocks til permit is available; Each release adds a permit
- Just keeps a count of available permits

- Semaphore constructor also accepts a fairness parameter:

*Semaphore(int permits, boolean fair);*

- Creates a Semaphore with the given number of permits and the given fairness setting.

*permits*: initial value

*fair*: if true semaphore uses FIFO to manage blocked threads; if set false, class doesn't guarantee order threads acquire permits.

- See `SemaApp.java` example in Labs

# High Level Concurrency Objects

- Previous examples
  - Low-level APIs that have been part of the Java platform from the very beginning.
  - Important to understand
  - To build simple applications
- Higher-level building blocks are needed for more advanced tasks, such as
  - massively concurrent applications that fully exploit today's multiprocessor and multi-core systems.
- High-level concurrency features introduced with version 5.0 of the Java platform
  - Implemented in the `java.util.concurrent` packages
  - New concurrent data structures in the Java Collections Framework.



# High Level Concurrency Objects:

## *java.util.concurrent*

- Lock objects
  - support locking idioms that simplify many concurrent applications.
- Executors
  - define a high-level API for launching and managing threads.
  - Executor implementations provided by *java.util.concurrent* provide thread pool management suitable for large-scale applications.
- Concurrent collections
  - make it easier to manage large collections of data, and can greatly reduce the need for synchronization.
- Atomic variables
  - have features that minimize synchronization and help avoid memory consistency errors.
- `ThreadLocalRandom`
  - for applications that expect to use random numbers from multiple threads

# Interface Lock

- Lock implementations provide more extensive locking operations than can be obtained using synchronized methods and statements.
- Their biggest advantage over implicit locks is can back out of an attempt to acquire a Lock:
  - i.e. livelock, starvation & deadlock are not a problem
- Lock methods:
  - *tryLock()* returns if lock is not available immediately or before a timeout (optional parameter) expires.
  - *lockInterruptibly()* returns if another thread sends an interrupt before the lock is acquired.

# Interface Condition

- *Conditions* (also known as condition queues or condition variables) provide
  - a means for one thread to suspend execution (to "wait") until notified by another thread that some state condition may now be true.
- Because access to this shared state information occurs in different threads, it must be protected, so a lock of some form is associated with the condition.
- The key property that waiting for a condition provides is:
  - it atomically releases the associated lock and suspends the current thread, just like *Object.wait*.

# Interface Lock

- Lock interface also supports a *wait/notify* mechanism, through the associated Condition objects
- Thus Lock and Condition replace basic monitor methods (`wait()`, `notify()` and `notifyAll()`) with specific objects:
  - Lock in place of *synchronized* methods and statements.
  - An associated Condition in place of Object's monitor methods.
  - A Condition instance is intrinsically bound to a Lock.
- To obtain a Condition instance for a particular Lock instance use its `newCondition()` method.

# Executors

- In the previous examples, there's a close connection between the task being done by a new thread, as defined by its Runnable object, and the thread itself, as defined by a Thread object.
- In large-scale applications, it makes sense to separate thread management and creation from the rest of the application.
  - Objects that encapsulate these functions are known as *executors*.

# Threads on a Multicore Machine

- How to execute the threads to different cores on a multicore machine?
- There are 2 mechanisms in Java
  - *Executor* Interface and Thread Pools
    - Executor – an interface that supports launching new tasks
    - Thread Pools - Address two different problems:
      - they usually provide improved performance when executing large numbers of asynchronous tasks;
      - they provide a means of bounding and managing the resources, including threads, consumed when executing a collection of tasks.
  - Fork/Join Framework

# Reference

- Chapter 3, *Introduction to Network Programming in Java* by Jan Graba
- *Parallel Programming in Java: A Tutorial*, Arik Baratz, Dror Birkman, Ofir Carny, Shy Cohen, and Assaf Schuster
- Java Tutorial section on Threads
  - <https://docs.oracle.com/javase/tutorial/essential/concurrency>
  - <https://docs.oracle.com/javase/tutorial/essential/concurrency/highlevel.html>