

CMPU4021 Distributed Systems

Middleware - Overview

Middleware


- Commonly heard term
- But no generally agreed meaning
- Most often used in the context of client server systems
 - Clients and servers communicate
 - This communicate often passes through intermediate software layers
 - This intermediate software is the middleware

Middleware: Attributes

- Provides services to applications
- Requires system resources, dependencies
- Has vulnerabilities and constraints
- May or may not implement its own access control model
- Developer may not have control over its design

Distributed Systems Middleware

- Enables the modular interconnection of distributed software
 - Usually via services



Common Middleware Services

- Naming and Directory Service
- Event Service
- Transaction Service
- Fault Detection Service
- Trading Service
- Replication Service
- Migration Service

Middleware dimensions

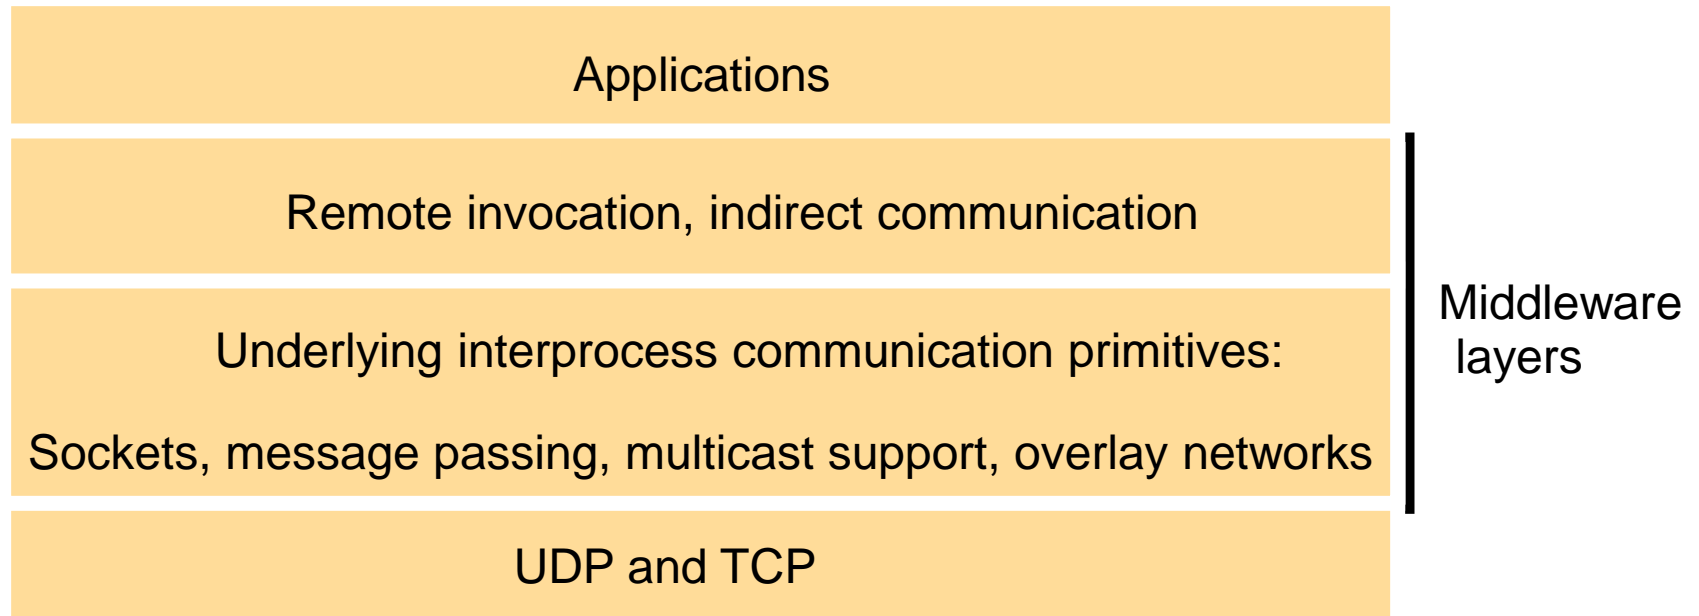
- Request/Reply Messaging vs. Asynchronous
- Language-specific vs. Language-independent
- Proprietary vs. Standards-based
- Small-scale vs. Large-scale
- Tightly-coupled components vs. Loosely-coupled

Middleware

- A whole range of technologies can be classified as middleware
- These include
 - RPC
 - Distributed objects and remote method invocation (*Java RMI, Corba*)
 - Remote event notification (*Jini*)
 - Remote SQL access (*JDBC*)
 - Distributed transaction processing

Middleware

- A suite of API software that uses underlying processes and communication (message passing) protocols to provide its abstract protocol – simple RMI request-reply protocol



Middleware

- Supports interaction between clients and servers
- Masks heterogeneity
- Provides a consistent programming model

Middleware layers

Middleware provides:

- *location transparency*:
 - RPC – the client that calls a procedure cannot tell whether the procedure runs in the same process or in a different process, different computer.
 - RMI – object making the invocation cannot tell whether the object it invokes is local or not;
 - EBP – the generating/receiving – not aware of one another's locations
- *protocol abstraction*
 - independent of underlying transport protocols

Middleware layers

Provides:

- *OS heterogeneity*
 - independent of the underlying operating system
- *hardware independence*
 - approaches to external data representations hide the differences due to hardware architectures, such as byte ordering.
- *multi-language support*
 - allows clients written in one language to invoke methods in objects that live in server programs written in another language.
 - Achieved by using an interface definition language (IDL) to define interfaces.

Middleware

- Middleware provides support for (some of):
 - Naming, Location, Service discovery, Replication
 - Protocol handling, Communication faults, QoS
 - Synchronisation, Concurrency, Transactions, Storage
 - Access control, Authentication

Asynchronous Middleware


- The client is not assumed to wait for the server after issuing request
- It may continue processing before reply arrives
- Often handled using message passing
 - hence, Message Oriented Middleware (MOM)

Persistent vs. Transient

- Another classification of communication
 - including middleware
- **persistent**: message life does not depend on continued sender execution
- **transient**: message life does depend on continued sender execution

Example

- Electronic mail
 - user message sent to local mail server
 - stored in temporary buffer
 - subsequently sent to target mail server
 - placed in mail box for recipient to read
 - note the places at which communication can be delayed waiting for delivery



Persistent Communication

- Message stored by communication system as long as it takes to deliver
- Sending application does not need to keep executing after sending
- Receiving applications does not have to be executing when message sent

Transient Communication

- Message stored only as long as both sending and receiving application are executing
- Can have transient asynchronous
 - both active
 - but sender continues immediately

Failure


- Persistent communication better at handling failures
 - network failure not a problem
 - other failures can be handled by retry (maybe)

Middleware - Types

- Remote Procedure Call (RPC)
- Object-Oriented Middleware (OOM)
 - Java RMI
 - CORBA
 - Reflective Middleware
 - Flexible middleware (OOM) for mobile and context-aware applications – adaptation to context
- Message-Oriented Middleware (MOM)
 - Java Message Service
 - IBM MQ
 - Web Services
- Event-Based Middleware
 - Jini

Object-Oriented Middleware (OOM)

- **Objects** can be *local* or *remote*
- **Object references** can be *local* or *remote*
- Remote objects have visible **remote interfaces**
- Masks remote objects as being local using **proxy objects**
- **Remote method invocation**



Properties of OOM


- ✓ Support for object-oriented programming model
 - objects, methods, interfaces, encapsulation, ...
 - exceptions (also in some RPC systems)
- ✓ **Location Transparency**
 - mapping object references to locations
- ✓ Synchronous request/reply interaction
 - same as RPC
- ✓ Services comprising multiple servers are easier to build with OOM

Disadvantages of OOM

- ✖ Synchronous request/reply interaction
 - Asynchronous Method Invocation (AMI)
 - But implementations may not be loosely coupled
- ✖ Distributed garbage collection
 - Releasing memory for unused remote objects
- ✖ OOM rather static and heavy-weight
 - Bad for ubiquitous systems and embedded devices

Event-Based Middleware/ Publish/Subscribe

- **Publishers** *publish events* (messages)
- **Subscribers** express interest in events with *subscriptions*
- **Event Service** *notifies* interested subscribers of published events
- Events can have arbitrary content or name/value pairs




Message Queues aka Message Oriented Middleware (MOM)

- Based on message passing
- Extensive support for persistent asynchronous communication
 - have intermediate-term storage capacity for messages
 - neither sender nor receiver required to be active during transmission
- Not a new idea
 - it is how networks work
 - for example, Unix sockets
- Messages can be large
 - time in minutes
 - as opposed to sockets, where seconds

Message-Oriented Middleware (MOM)

- Communication using **messages**
- Messages stored in **message queues**
- Optional **message server** decouples client and server
- Various assumptions about **message content**



Properties of MOM

- ✓ **Asynchronous** interaction
 - Client and server are only **loosely coupled**
 - Messages are queued
 - Good for application integration
- ✓ Support for **reliable** delivery service
 - Keep queues in persistent storage
- ✓ Processing of messages by intermediate message server
 - Filtering, transforming, logging, ...
 - Networks of message servers
- ✓ Natural for database integration

Disadvantages of MOM

- ✘ Poor programming abstracting
 - Rather low-level (cf. Packets)
 - Results in multi-threaded code
 - Request/reply more difficult to achieve
- ✘ Message formats unknown to middleware
 - No type checking
- ✘ Queue abstraction only gives one-to-one communication
 - Limits scalability

MOM/MQ - additional functionalities

- Transactions support
 - Support for the sending or receiving of a message to be contained within a transaction
- Message transformation
 - An arbitrary transformation can be performed on an arriving message.
 - E.g. - to transform messages between formats to deal with heterogeneity in underlying data representations.
 - Important tool in dealing with heterogeneity
 - Message broker
 - Term often used to denote a service responsible for message transformation.

Message queues vs Message Passing

- Message queues are similar to the message-passing systems
- The difference
 - message-passing systems have *implicit* queues associated with senders and receivers
 - message queuing systems have *explicit* queues that are third-party entities, *separate* from the sender and the receiver.
- This is the key difference that makes message queues
 - an indirect communication paradigm
 - with the crucial properties
 - of space and time uncoupling.

MOM Examples/Toolkits

- A major class of commercial middleware with key implementations including
 - IBM's MQ (previously WebSphere MQ),
 - Microsoft's MSMQ
 - Oracle's Streams Advanced Queuing (AQ).
- Other Examples:
 - Java Message Service
 - Web Services
- The MOM paradigm has had a long history in distributed applications.
 - Message Queue Services (MQS) have been in use since the 1980's.

IBM MQ

- IBM messaging middleware.
- Provides messaging and queuing capabilities across multiple modes of operation:
 - point-to-point
 - publish/subscribe

IBM MQ

Provides:

- Messaging
 - Programs communicate by sending each other data in messages rather than by calling each other directly.
- Queuing
 - Messages are placed on queues, so that programs can run independently of each other, at different speeds and times, in different locations, and without having a direct connection between them.
- Point-to-point
 - Applications send messages to a queue and receive messages from a queue. Each message is consumed by a single instance of an application. The sender must know the name of the destination, but not where it is.
- Publish/subscribe
 - Applications subscribe to topics. When an application publishes a message on a topic, IBM MQ sends copies of the message to those subscribing applications. The publisher does not know the names of subscribers, or where they are.

Java Message Service (JMS)

- A specification of a standardized way for distributed Java programs to communicate indirectly
 - API specification to access MOM implementations
- Two modes of operation:
 - Point-to-point
 - one-to-one communication using queues
 - Publish/Subscribe
 - As an Event-Based Middleware
- Implementations
 - Joram from OW2,
 - Apache ActiveMQ
 - IBM MQ, also provide a JMS interface on to their underlying infrastructure.

Key roles in JMS

- A JMS client
 - a Java program or component that produces or consumes messages
- JMS producer
 - a program that creates and produces messages
- JMS consumer
 - a program that receives and consumes messages.
- JMS provider
 - any of the multiple systems that implement the JMS specification.
- JMS message
 - an object that is used to communicate information between JMS clients (from producers to consumers).
 - Java objects can be serialised to JMS messages

Criteria for selecting *middleware*

- Suitability
 - integration of software/hardware aspects of architectures
 - Users will only be satisfied if their middleware–OS combination has good performance.
 - Middleware runs on a variety of OS–hardware combinations (platforms) at the nodes of a distributed system.
- Integration of applications
 - standards and middleware technology considerations
- Reliability and robustness
- Transparency
- Risks and cost aspects

Criteria for selecting *middleware*

- Strength of product support
 - The maturity and stability of the tool;
 - The fault tolerance provided by the tool;
 - The availability of developer tools;
 - Maintainability;
 - Code reuse
- Security characteristics

Middleware: Security Goals

- Engineer application to depend on middleware only as much as necessary, in view of middleware's capabilities, liabilities and constraints
- Engineer system to account for middleware's capabilities, liabilities and constraints.

References

- Coulouris, Dollimore and Kindberg, Distributed Systems: Concepts and Design, Section 5.1
- Dr J. Bacon, University of Cambridge